

bm-project

YISU

2024-12-19

```
# Load necessary libraries
library(tidyverse)
library(broom)
library(GGally)

## Registered S3 method overwritten by 'GGally':
##   method from
##   +.gg      ggplot2

library(gridExtra)

##
## Attaching package: 'gridExtra'

## The following object is masked from 'package:dplyr':
##
##   combine

# Read the dataset
data <- read.csv("Project_1_data.csv")

# Inspect the structure of the data
glimpse(data)

## Rows: 948
## Columns: 14
## $ Gender          <chr> "female", "female", "female", "male", "male", "fem~
## $ EthnicGroup     <chr> "", "group C", "group B", "group A", "group C", "g~
## $ ParentEduc      <chr> "bachelor's degree", "some college", "master's deg~
## $ LunchType       <chr> "standard", "standard", "standard", "free/reduced"~
## $ TestPrep        <chr> "none", "", "none", "none", "none", "none", "compl~
## $ ParentMaritalStatus <chr> "married", "married", "single", "married", "marrie~
## $ PracticeSport   <chr> "regularly", "sometimes", "sometimes", "never", "s~
## $ IsFirstChild    <chr> "yes", "yes", "yes", "no", "yes", "yes", "no", "ye~
## $ NrSiblings      <int> 3, 0, 4, 1, 0, 1, 1, 1, 3, NA, 1, 1, 1, 1, 2, 0, 0~
## $ TransportMeans  <chr> "school_bus", "", "school_bus", "", "school_bus", ~
## $ WklyStudyHours  <chr> "< 5", "10-May", "< 5", "10-May", "10-May", "10-Ma~
## $ MathScore       <int> 71, 69, 87, 45, 76, 73, 85, 41, 65, 37, 58, 40, 66~
## $ ReadingScore    <int> 71, 90, 93, 56, 78, 84, 93, 43, 64, 59, 54, 52, 82~
## $ WritingScore    <int> 74, 88, 91, 42, 75, 79, 89, 39, 68, 50, 52, 43, 74~

hist_math <- ggplot(data, aes(x = MathScore)) +
  geom_histogram(binwidth = 5, fill = "steelblue", color = "black") +
  labs(title = "Histogram of Math Scores", x = "Math Score", y = "Frequency")
```

```

hist_reading <- ggplot(data, aes(x = ReadingScore)) +
  geom_histogram(binwidth = 5, fill = "firebrick", color = "black") +
  labs(title = "Histogram of Reading Scores", x = "Reading Score", y = "Frequency")

hist_writing <- ggplot(data, aes(x = WritingScore)) +
  geom_histogram(binwidth = 5, fill = "darkgreen", color = "black") +
  labs(title = "Histogram of Writing Scores", x = "Writing Score", y = "Frequency")
box_nr_siblings <- ggplot(data, aes(y = NrSiblings)) +
  geom_boxplot(fill = "lightblue") +
  labs(title = "Boxplot of Number of Siblings", y = "Number of Siblings")

box_wkly_study_hours <- ggplot(data, aes(y = WklyStudyHours)) +
  geom_boxplot(fill = "lightgreen") +
  labs(title = "Boxplot of Weekly Study Hours", y = "Weekly Study Hours")
scatter_math_reading <- ggplot(data, aes(x = ReadingScore, y = MathScore)) +
  geom_point(color = "darkblue") +
  labs(title = "Math Score vs. Reading Score", x = "Reading Score", y = "Math Score")

scatter_reading_writing <- ggplot(data, aes(x = WritingScore, y = ReadingScore)) +
  geom_point(color = "darkred") +
  labs(title = "Reading Score vs. Writing Score", x = "Writing Score", y = "Reading Score")
# Convert categorical variables to factors if needed
data <- data %>%
  mutate(across(c(Gender, EthnicGroup, ParentEduc, LunchType, TestPrep,
    ParentMaritalStatus, PracticeSport, IsFirstChild,
    TransportMeans, WklyStudyHours), as.factor))

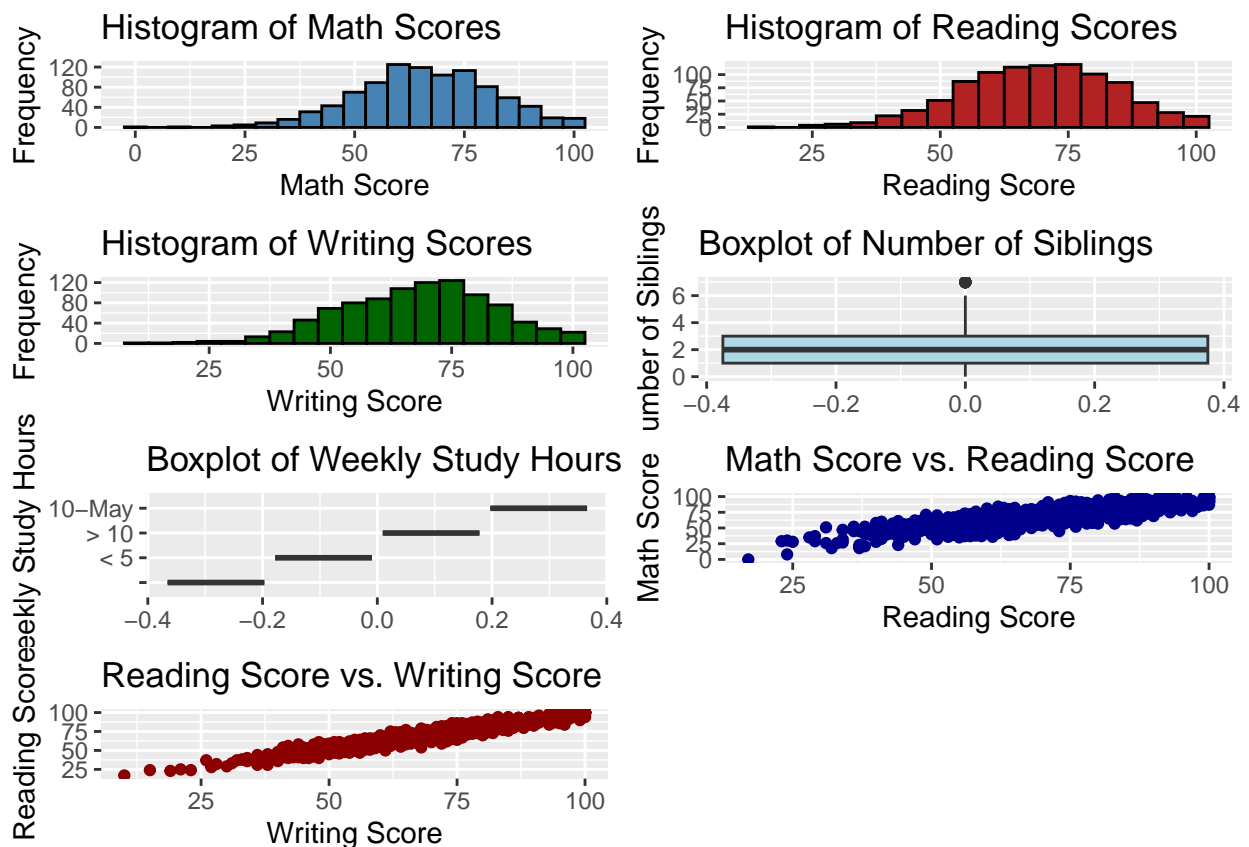
pairwise_plots <- ggpairs(data,
  columns = c("MathScore", "ReadingScore", "WritingScore", "NrSiblings", "WklyStudyHours"),
  aes(color = Gender),
  lower = list(continuous = "smooth"),
  upper = list(continuous = "cor"),
  diag = list(continuous = "densityDiag"))
grid.arrange(
  hist_math, hist_reading, hist_writing,
  box_nr_siblings, box_wkly_study_hours,
  scatter_math_reading, scatter_reading_writing,
  ncol = 2
)

```

```

## Warning: Removed 46 rows containing non-finite outside the scale range
## (`stat_boxplot()`).

```



```
data$Gender <- as.factor(data$Gender)
data$EthnicGroup <- as.factor(data$EthnicGroup)
data$ParentEduc <- as.factor(data$ParentEduc)
data$LunchType <- as.factor(data$LunchType)
data$TestPrep <- as.factor(data$TestPrep)
data$ParentMaritalStatus <- as.factor(data$ParentMaritalStatus)
data$PracticeSport <- as.factor(data$PracticeSport)
data$IsFirstChild <- as.factor(data$IsFirstChild)
data$TransportMeans <- as.factor(data$TransportMeans)
data$WklyStudyHours <- as.factor(data$WklyStudyHours)

# Handle missing values by removing rows with NA
data <- na.omit(data)

# Define the full model for MathScore
math_full_model <- lm(MathScore ~ Gender + EthnicGroup + ParentEduc + LunchType + TestPrep +
  ParentMaritalStatus + PracticeSport + IsFirstChild + NrSiblings +
  TransportMeans + WklyStudyHours, data = data)

# Perform stepwise model selection
math_selected_model <- stepAIC(math_full_model, direction = "both", trace = FALSE)
math_selected_model
```

```
##
## Call:
## lm(formula = MathScore ~ Gender + EthnicGroup + ParentEduc +
##   LunchType + TestPrep + ParentMaritalStatus + IsFirstChild +
```

```

##      WklyStudyHours, data = data)
##
## Coefficients:
##              (Intercept)              Gendermale
##              51.2910              5.0885
##      EthnicGroupgroup A      EthnicGroupgroup B
##              -1.2758              0.1118
##      EthnicGroupgroup C      EthnicGroupgroup D
##              -0.2774              3.5351
##      EthnicGroupgroup E ParentEducassociate's degree
##              8.5754              4.7280
##      ParentEducbachelor's degree      ParentEduchigh school
##              6.0476              -0.8361
##      ParentEducmaster's degree      ParentEducsome college
##              6.4044              3.7222
##      ParentEducsome high school      LunchTypestandard
##              -0.5013              11.0841
##      TestPrepcompleted      TestPrepnone
##              4.4017              -0.9571
##      ParentMaritalStatusdivorced      ParentMaritalStatusmarried
##              -0.5039              3.2310
##      ParentMaritalStatussingle      ParentMaritalStatuswidowed
##              0.1923              4.2876
##      IsFirstChildno      IsFirstChildyes
##              -1.1316              0.8604
##      WklyStudyHours< 5      WklyStudyHours> 10
##              -4.7980              -1.0807
##      WklyStudyHours10-May
##              -1.4518

# Display the summary of the selected model
cat("\nSelected Model for MathScore:\n")

##
## Selected Model for MathScore:
summary(math_selected_model)

##
## Call:
## lm(formula = MathScore ~ Gender + EthnicGroup + ParentEduc +
##      LunchType + TestPrep + ParentMaritalStatus + IsFirstChild +
##      WklyStudyHours, data = data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -49.527  -8.997   0.498   9.626  30.285
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      51.2910     5.1078  10.042 < 2e-16 ***
## Gendermale        5.0885     0.8992   5.659 2.06e-08 ***
## EthnicGroupgroup A -1.2758     2.4139  -0.529 0.597264
## EthnicGroupgroup B  0.1118     2.0927   0.053 0.957422
## EthnicGroupgroup C -0.2774     1.9986  -0.139 0.889628
## EthnicGroupgroup D  3.5351     2.0268   1.744 0.081489 .

```

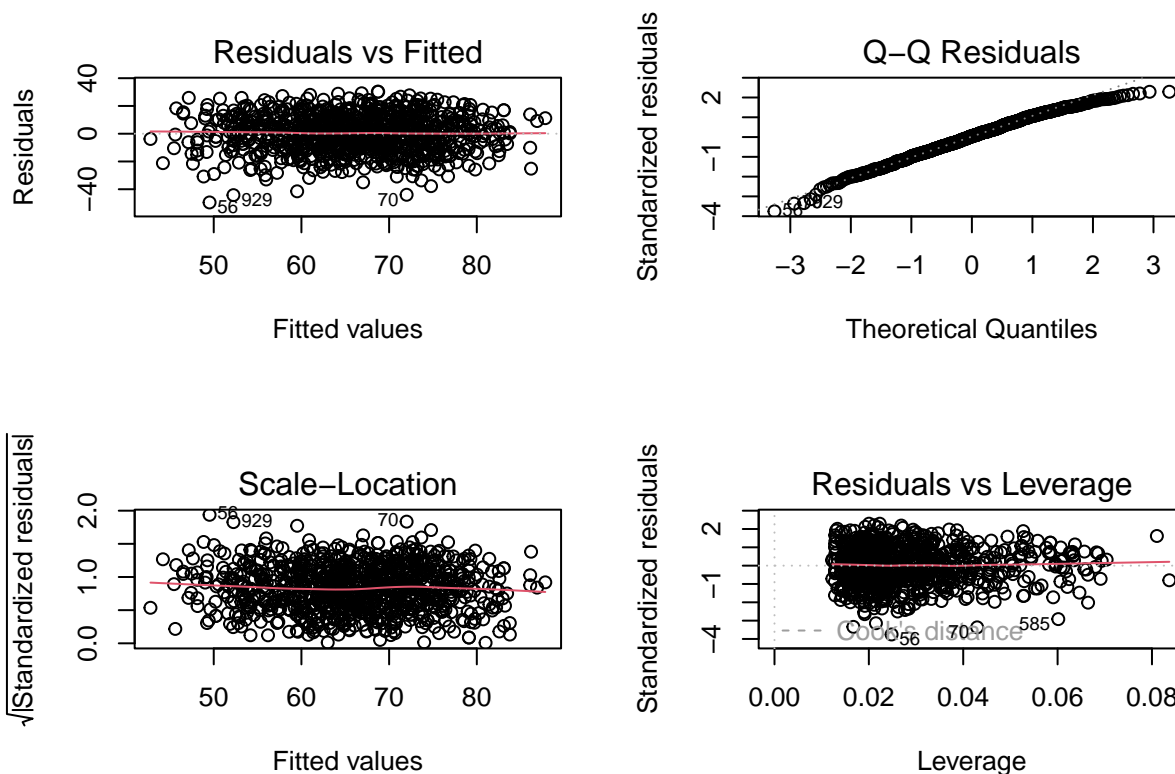
```
## EthnicGroupgroup E      8.5754      2.2003      3.897 0.000105 ***
## ParentEducassociate's degree  4.7280      2.1442      2.205 0.027708 *
## ParentEducbachelor's degree  6.0476      2.3515      2.572 0.010280 *
## ParentEduchigh school -0.8361      2.1666     -0.386 0.699661
## ParentEducmaster's degree  6.4044      2.6695      2.399 0.016643 *
## ParentEducsome college  3.7222      2.1417      1.738 0.082561 .
## ParentEducsome high school -0.5013      2.1831     -0.230 0.818444
## LunchTypestandard      11.0841      0.9381     11.815 < 2e-16 ***
## TestPrepcompleted      4.4017      2.0171      2.182 0.029357 *
## TestPrepnone -0.9571      1.9433     -0.493 0.622484
## ParentMaritalStatusdivorced -0.5039      2.3146     -0.218 0.827721
## ParentMaritalStatusmarried  3.2310      2.0944      1.543 0.123277
## ParentMaritalStatussingle  0.1923      2.2226      0.087 0.931062
## ParentMaritalStatuswidowed  4.2876      3.4046      1.259 0.208240
## IsFirstChildno -1.1316      2.6597     -0.425 0.670602
## IsFirstChildyes  0.8604      2.6059      0.330 0.741341
## WklyStudyHours< 5 -4.7980      2.3751     -2.020 0.043676 *
## WklyStudyHours> 10 -1.0807      2.4903     -0.434 0.664422
## WklyStudyHours10-May -1.4518      2.3020     -0.631 0.528423
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 13.36 on 877 degrees of freedom
## Multiple R-squared:  0.2819, Adjusted R-squared:  0.2623
## F-statistic: 14.35 on 24 and 877 DF, p-value: < 2.2e-16

# Make predictions using the selected model
math_predictions <- predict(math_selected_model, newdata = data)

# Combine actual and predicted values for MathScore
math_results <- data.frame(
  Actual_MathScore = data$MathScore,
  Predicted_MathScore = math_predictions
)

# Save results to a CSV file
write.csv(math_results, "math_score_predictions.csv", row.names = FALSE)

par(mfrow = c(2, 2))
plot(math_selected_model)
```



```
par(mfrow = c(1, 1))

reading_full_model <- lm(ReadingScore ~ Gender + EthnicGroup + ParentEduc + LunchType + TestPrep +
  ParentMaritalStatus + PracticeSport + IsFirstChild + NrSiblings +
  TransportMeans + WklyStudyHours, data = data)

# Perform stepwise model selection
reading_selected_model <- stepAIC(reading_full_model, direction = "both", trace = FALSE)
reading_selected_model

##
## Call:
## lm(formula = ReadingScore ~ Gender + EthnicGroup + ParentEduc +
##   LunchType + TestPrep + ParentMaritalStatus + IsFirstChild,
##   data = data)
##
## Coefficients:
##              (Intercept)              Gendermale
##              58.63929              -7.15237
##   EthnicGroupgroup A      EthnicGroupgroup B
##              2.11437              1.61464
##   EthnicGroupgroup C      EthnicGroupgroup D
##              2.10744              4.82071
##   EthnicGroupgroup E ParentEducassociate's degree
##              6.46862              5.14291
## ParentEducbachelor's degree      ParentEduchigh school
##              6.13542              -0.87311
##   ParentEducmaster's degree      ParentEducsome college
##              9.02888              3.27435
## ParentEducsome high school      LunchTypestandard
```

```
##              0.05839              7.56883
##      TestPrepcompleted      TestPrepnone
##      6.15804      -0.75034
## ParentMaritalStatusdivorced ParentMaritalStatusmarried
##      -1.08867      2.67064
## ParentMaritalStatussingle ParentMaritalStatuswidowed
##      -0.53098      3.47406
##      IsFirstChildno      IsFirstChildyes
##      -1.44383      0.71327

# Display the summary of the selected model
cat("\nSelected Model for Reading:\n")

##
## Selected Model for Reading:

summary(reading_selected_model)

##
## Call:
## lm(formula = ReadingScore ~ Gender + EthnicGroup + ParentEduc +
##      LunchType + TestPrep + ParentMaritalStatus + IsFirstChild,
##      data = data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -43.237  -8.650   1.161   9.385  29.517
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      58.63929      4.50285    13.023 < 2e-16 ***
## Gendermale       -7.15237      0.88045   -8.124 1.53e-15 ***
## EthnicGroupgroup A      2.11437      2.35196    0.899 0.368908
## EthnicGroupgroup B      1.61464      2.04598    0.789 0.430219
## EthnicGroupgroup C      2.10744      1.94808    1.082 0.279636
## EthnicGroupgroup D      4.82071      1.97981    2.435 0.015092 *
## EthnicGroupgroup E      6.46862      2.15021    3.008 0.002701 **
## ParentEducassociate's degree  5.14291      2.09638    2.453 0.014351 *
## ParentEducbachelor's degree  6.13542      2.29642    2.672 0.007685 **
## ParentEduchigh school   -0.87311      2.11963   -0.412 0.680499
## ParentEducmaster's degree  9.02888      2.60907    3.461 0.000565 ***
## ParentEducsome college   3.27435      2.09170    1.565 0.117848
## ParentEducsome high school  0.05839      2.13514    0.027 0.978191
## LunchTypestandard      7.56883      0.91763    8.248 5.83e-16 ***
## TestPrepcompleted      6.15804      1.96245    3.138 0.001758 **
## TestPrepnone          -0.75034      1.89705   -0.396 0.692548
## ParentMaritalStatusdivorced -1.08867      2.26625   -0.480 0.631072
## ParentMaritalStatusmarried  2.67064      2.05082    1.302 0.193177
## ParentMaritalStatussingle -0.53098      2.17544   -0.244 0.807227
## ParentMaritalStatuswidowed  3.47406      3.33026    1.043 0.297151
## IsFirstChildno       -1.44383      2.60496   -0.554 0.579541
## IsFirstChildyes        0.71327      2.55239    0.279 0.779964
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
```

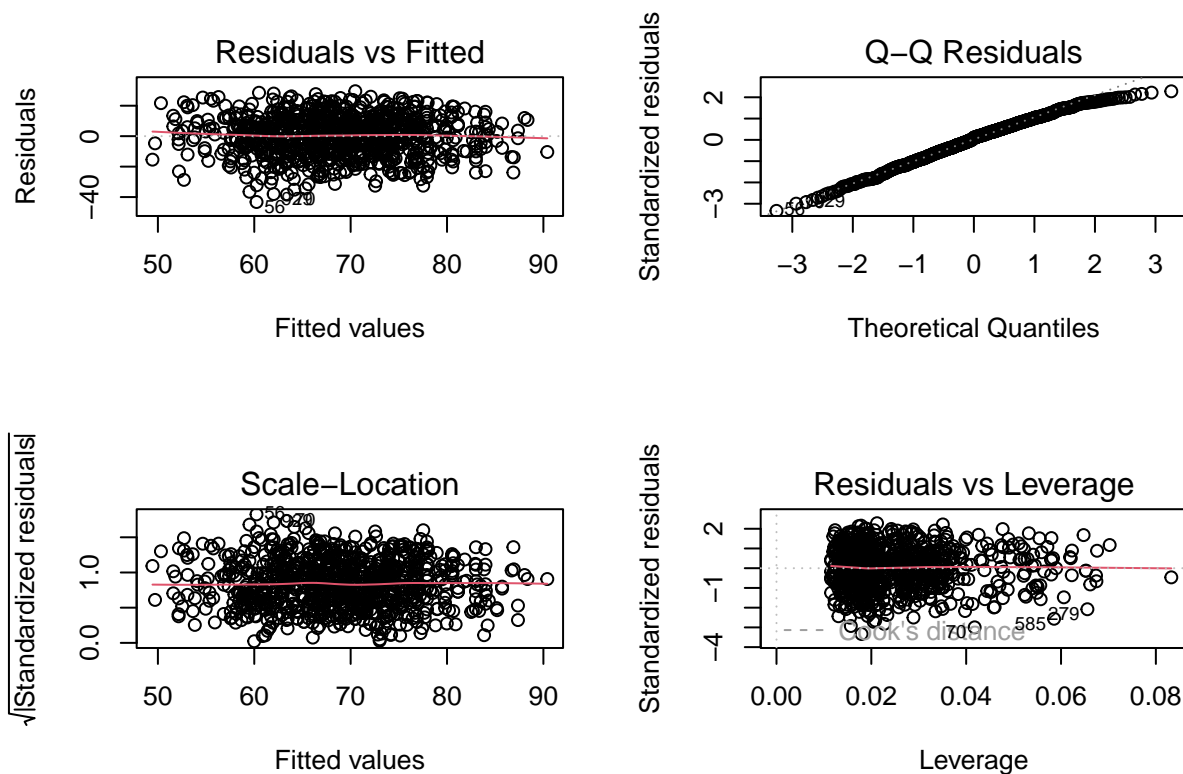
```
## Residual standard error: 13.09 on 880 degrees of freedom
## Multiple R-squared:  0.2389, Adjusted R-squared:  0.2207
## F-statistic: 13.15 on 21 and 880 DF,  p-value: < 2.2e-16

# Make predictions using the selected model
reading_predictions <- predict(reading_selected_model, newdata = data)

# Combine actual and predicted values for MathScore
reading_results <- data.frame(
  Actual_ReadingScore = data$ReadingScore,
  Predicted_ReadingScore = reading_predictions
)

# Save results to a CSV file
write.csv(reading_results, "reading_score_predictions.csv", row.names = FALSE)

par(mfrow = c(2, 2))
plot(reading_selected_model)
```



```
par(mfrow = c(1, 1))

writing_full_model <- lm(WritingScore ~ Gender + EthnicGroup + ParentEduc + LunchType + TestPrep +
  ParentMaritalStatus + PracticeSport + IsFirstChild + NrSiblings +
  TransportMeans + WklyStudyHours, data = data)

# Perform stepwise model selection
writing_selected_model <- stepAIC(writing_full_model, direction = "both", trace = FALSE)
writing_selected_model

##
## Call:
```



```
## lm(formula = WritingScore ~ Gender + EthnicGroup + ParentEduc +
##     LunchType + TestPrep + ParentMaritalStatus + WklyStudyHours,
##     data = data)
##
## Coefficients:
##             (Intercept)                Gendermale
##             59.36546                -9.23816
##           EthnicGroupgroup A           EthnicGroupgroup B
##             0.63448                0.66800
##           EthnicGroupgroup C           EthnicGroupgroup D
##             1.27259                5.76199
##           EthnicGroupgroup E ParentEducassociate's degree
##             5.03801                4.48934
## ParentEducbachelor's degree           ParentEduchigh school
##             6.69286                -2.30978
## ParentEducmaster's degree           ParentEducsome college
##             9.76964                3.19346
## ParentEducsome high school           LunchTypestandard
##             -1.80483                8.31180
##           TestPrepcompleted           TestPrepnone
##             7.72567                -1.82437
## ParentMaritalStatusdivorced           ParentMaritalStatusmarried
##             -0.34162                3.20838
## ParentMaritalStatussingle           ParentMaritalStatuswidowed
##             -0.05569                3.15168
##           WklyStudyHours< 5           WklyStudyHours> 10
##             -2.17429                -0.16869
##           WklyStudyHours10-May
##             0.34306

# Display the summary of the selected model
cat("\nSelected Model for WritingScore:\n")

##
## Selected Model for WritingScore:
summary(writing_selected_model)

##
## Call:
## lm(formula = WritingScore ~ Gender + EthnicGroup + ParentEduc +
##     LunchType + TestPrep + ParentMaritalStatus + WklyStudyHours,
##     data = data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -46.784  -7.859   0.772   8.712  32.001
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    59.36546    4.20505  14.118 < 2e-16 ***
## Gendermale     -9.23816    0.85274 -10.834 < 2e-16 ***
## EthnicGroupgroup A    0.63448    2.28724   0.277 0.781539
## EthnicGroupgroup B    0.66800    1.98226   0.337 0.736205
## EthnicGroupgroup C    1.27259    1.89387   0.672 0.501790
```

```

## EthnicGroupgroup D          5.76199    1.92143    2.999 0.002787 **
## EthnicGroupgroup E          5.03801    2.08690    2.414 0.015977 *
## ParentEducassociate's degree 4.48934    2.03180    2.210 0.027394 *
## ParentEducbachelor's degree 6.69286    2.22839    3.003 0.002745 **
## ParentEduchigh school      -2.30978    2.05029   -1.127 0.260235
## ParentEducmaster's degree   9.76964    2.52630    3.867 0.000118 ***
## ParentEducsome college      3.19346    2.02703    1.575 0.115516
## ParentEducsome high school -1.80483    2.06990   -0.872 0.383478
## LunchTypestandard           8.31180    0.88936    9.346 < 2e-16 ***
## TestPrepcompleted           7.72567    1.91429    4.036 5.92e-05 ***
## TestPrepnone               -1.82437    1.84365   -0.990 0.322669
## ParentMaritalStatusdivorced -0.34162    2.19391   -0.156 0.876295
## ParentMaritalStatusmarried  3.20838    1.98472    1.617 0.106336
## ParentMaritalStatussingle  -0.05569    2.10631   -0.026 0.978913
## ParentMaritalStatuswidowed  3.15168    3.23112    0.975 0.329623
## WklyStudyHours< 5          -2.17429    2.25368   -0.965 0.334924
## WklyStudyHours> 10         -0.16869    2.36248   -0.071 0.943092
## WklyStudyHours10-May        0.34306    2.18355    0.157 0.875192
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 12.68 on 879 degrees of freedom
## Multiple R-squared:  0.3347, Adjusted R-squared:  0.3181
## F-statistic: 20.1 on 22 and 879 DF, p-value: < 2.2e-16

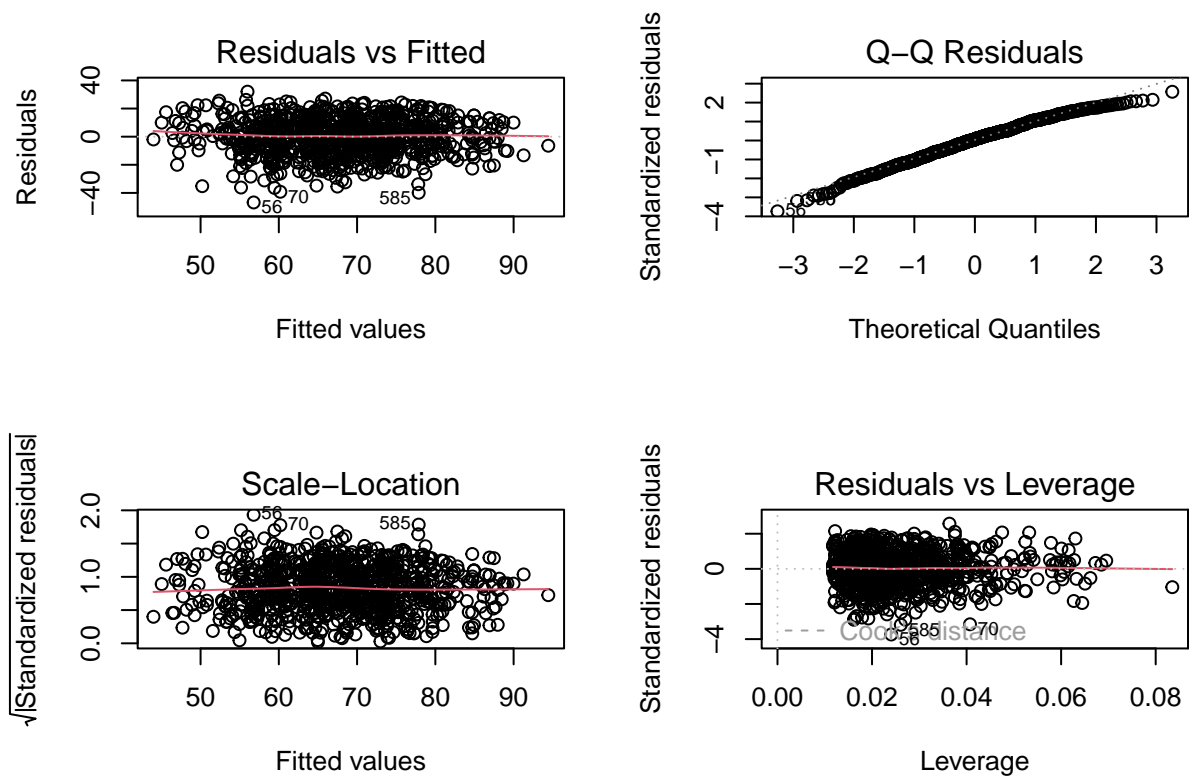
# Make predictions using the selected model
writing_predictions <- predict(writing_selected_model, newdata = data)

# Combine actual and predicted values for MathScore
writing_results <- data.frame(
  Actual_WritingScore = data$WritingScore,
  Predicted_WritingScore = writing_predictions
)

# Save results to a CSV file
write.csv(writing_results, "writing_score_predictions.csv", row.names = FALSE)

par(mfrow = c(2, 2))
plot(writing_selected_model)

```



```
par(mfrow = c(1, 1))
```